

Other Implementation Details

Model Configuration

1. Overcorrection Evaluation: over_correction_evl

Configuration files:
[model_groups.json](#)
[zero_shot_models.json](#)

The overcorrection evaluation uses 11 models organized into two benchmark groups: **Small** and **Large**. The Small group contains locally hosted models served through Ollama, while the Large group contains hosted API models. The All group combines both sets.

Table X. Model configurations used in the overcorrection evaluation.

Group	Model	Model Identifier	Runtime	Temperature	Max Tokens	Reasoning / Thinking
Small	Qwen3.5 9B	qwen3.5:9b	Ollama	0.0	256	Disabled
Small	GLM-4 9B	glm4:9b	Ollama	0.0	512	—
Small	DeepSeek-R1 8B	deepseek-r1:8b	Ollama	0.0	2,048	Enabled (Low)
Small	Llama 3.1 8B	llama3.1:latest	Ollama	0.0	512	—
Small	Mistral 7B	mistral:latest	Ollama	0.0	512	—
Small	Gemma 2 9B	gemma2:9b	Ollama	0.0	512	—
Large	DeepSeek V4 Pro	deepseek-v4-pro	API	0.0	512	Enabled
Large	Qwen 3.7 Plus	qwen3.7-plus	API	0.0	4,096	Enabled
Large	GLM-5.2	glm-5.2	API	0.0	4,096	Enabled
Large	GPT-4o	gpt-4o-2024-11-20	API	0.0	512	—
Large	Gemini 3.6 Flash	gemini-3.6-flash	API	—	512	Enabled (Low)

The All benchmark group contains all 11 models, while the OpenAI/Google group contains GPT-4o and Gemini 3.6 Flash. The complete model registry and group definitions are provided in the configuration files.

2. Intervention Gate and Rewriting: gender_intervention_gate

Configuration file:
[models.yaml](#)

The PAIFAR intervention and rewriting pipeline supports both hosted API models and locally hosted models. The configuration specifies the inference provider, model identifier, decoding parameters, and reasoning settings.

Table X. API and provider-level model configurations.

Model	Model Identifier	Runtime	Temperature	Max Tokens	Reasoning / Thinking
-------	------------------	---------	-------------	------------	----------------------

Model	Model Identifier	Runtime	Temperature	Max Tokens	Reasoning / Thinking
OpenAI	Environment-configured	API	—	—	—
GPT-4o Judge	gpt-4o	API	—	—	—
Gemini	Environment-configured	API	—	—	—
DeepSeek	Environment-configured	API	—	—	Enabled (High)
Qwen	Environment-configured	API	—	—	Enabled (Budget: 512)
GLM	Environment-configured	API	—	—	Enabled (High)
Local	Environment-configured	API	—	—	—

Table X. Local model configurations.

Model	Model Identifier	Runtime	Temperature	Max Tokens	Reasoning / Thinking
Qwen3.5 9B	qwen3.5:9b	Ollama	0.0	256	Disabled
DeepSeek LLM 7B	deepseek-llm:7b	Ollama	0.0	512	—
DeepSeek-R1 8B	deepseek-r1:8b	Ollama	0.0	2,048	Disabled
GLM-4 9B	glm4:9b	Ollama	0.0	512	—
Gemma 2 9B	gemma2:9b	Ollama	0.0	512	—
Llama 3.1 8B	llama3.1:latest	Ollama	0.0	128 / 256*	—
Mistral 7B	mistral:latest	Ollama	0.0	512	—

* Llama 3.1 8B uses a 128-token configuration for structured label prediction and a 256-token configuration for free-text rewriting.